

# ECG classification using deep neural networks: Investigating architecture optimization and transfer learning

**Pawel Kudzia**  
Simon Fraser University

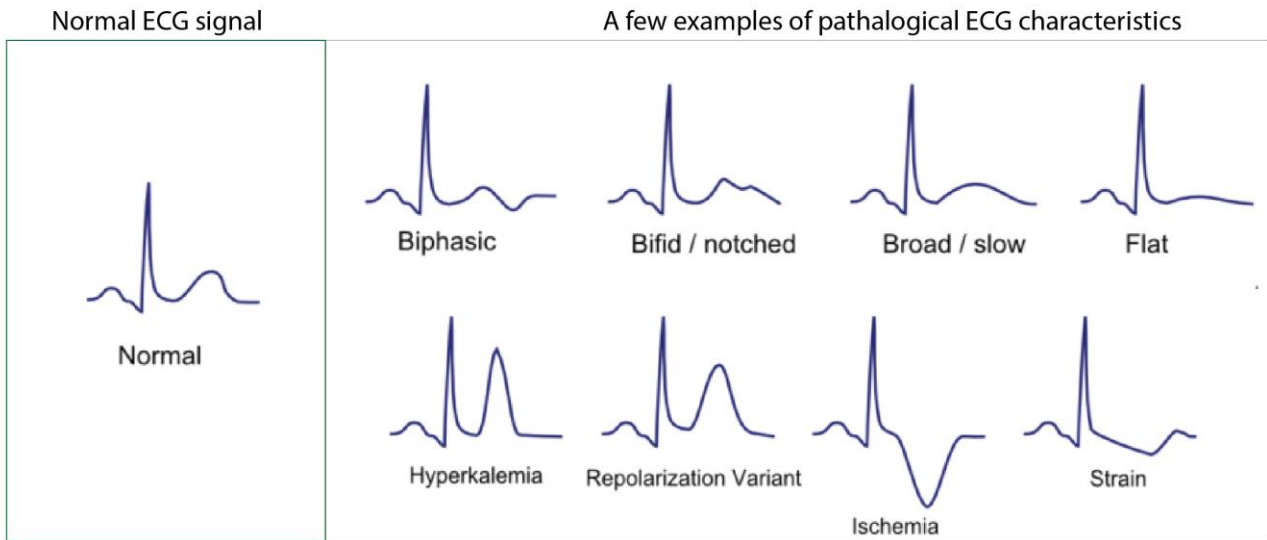
**Abstract** — Electrocardiogram (ECG) interpretation is an important medical diagnostic tool for diagnosing conditions of the heart. The use of neural networks aiming to assist cardiologists in the interpretation of ECG signals show great potential and could lead to improved patient outcomes. Here, I compared and developed several neural network architectures for classifying ECG signals from widely available datasets. Additionally, I compared modelling results when using augmented data and evaluated how well the models trained on one dataset classified data from another dataset. I found that the presented neural network architectures accurately classified pathological ECG signals, comparable to that of several published articles. I also found that data augmentation diminished overall performance, suggesting the approach taken was too simple. Finally, I found that transfer learning is possible but, classification performance was poor on pathological ECG.

**Keywords** —ECG, deep neural networks, architecture optimization, transfer learning, classification problems, error matrix

## 1.0 Introduction

An electrocardiogram (ECG) is a measurement of the electrical activity of the heart muscle. Each time the heart pumps blood it produces a characteristic time series waveform with distinct and identifiable features (right, Figure 1). In medicine, the ECG signal is used as a diagnostic tool to determine if the heart is functioning as expected. When someone has a pathology, such as heart valve leaks or tears, there can be subtle but detectable changes in the ECG signal (left, Figure 1). The ability to detect the presence of a pathology early could lead to quick diagnosis and improved outcomes for patients, not to mention reductions in long term healthcare costs [1]–[4].

Computer-aided interpretation of ECG signals has been an increasingly important tool in the clinical workflow, with first attempts dating back to the late 1950s [3]. Although sophisticated computer algorithms have shown promise, high rates of misdiagnosis are common (~50%). The presence of noise, variability in wave morphology between patients, and the mixture of both subjective and objective characteristics makes ECG classification a difficult problem, where even experienced cardiologists can disagree on prognosis [5]. The use of deep neural networks has led to major advances in image classification, speech recognition, and many medical diagnostics [1], [6], [7]. The ability of deep neural networks to recognize patterns and learn features of time varying signals without extensive preprocessing makes them particularly well suited for ECG classification.



**Figure 1:** Example characteristics of normal and pathological ECG signals.

There are many possible pathologies that can be observed in ECG waveforms. It is common for researchers to study only certain more common pathologies, or group pathologies into categories due to limitations in the amount of data on pathological ECG signals [1], [4], [8]. In general, many research groups have shown great success with ECG classification using neural networks, achieving low misclassification errors, and accuracy's nearing 99% [2], [7], [9]–[13]. Although the results are profound and have led to startup companies and major advances, the caveats are small representative test sets, with limited patients, and limited variations in the ECG pathologies themselves. This increases the chances of missing abnormalities that may exist warranting expert cardiologist option (who can also miss these suitable features). Recently, a monumental effort led by Andrew Ng's group at Stanford university, proposed a neural network architecture for evaluating up to 12 different pathologies (more then has ever been proposed before) [3], [7]. This effort was only possible due to the extensive amount of data the group was able to collect (over 50,000 patients with all ECG annotations performed by cardiologists) which made it possible to train and evaluate their model.

This project here aims to use deep neural networks to classify normal and pathological ECG waveforms from publically available datasets. Specifically, there are several objectives. The first objective is build, develop, and train a convolutional neural network using a previously established architecture and compare its ability to classify normal and 4 pathological ECG signals to the ability of two author developed architectures. The second objective is characterize the performance variations when a simple data augmentation is performed on the dataset. Finally, the third objective is to quantify the performance of the models established in aim 1 and 2 on a novel ECG dataset via transfer learning.

## 2.0 Methods

### 2.1 Data Curation

ECG data was taken from the MIT-BIH arrhythmia (48 subjects,  $f = 360$  Hz) and the PTB Diagnostic ECG (290 subjects,  $f = 1000$  Hz) datasets as made available through Kaggle courtesy of [1]. Both of these ECG datasets have been labelled by trained cardiologists. The dataset has some preprocessing steps that were taken which modify it from its raw time-series form. In brief, the continuous ECG voltage signal was divided into 10s windows, voltage normalized to be between 0-1, peaks were identified using the first derivative, and the median time interval between peaks was determined ( $T_{interval}$ ). Using the  $T_{interval}$  determined for each 10 second window, the signal was then divided up such that each frame starts from a peak and continues for  $1.2 * T_{interval}$ . Lastly, to ensure the divisions of each new frame have a fixed length, each frame is padded with zeros to make its length equal to an array of 187 in length. Each interval is then given the corresponding label which the cardiologist has previously provided. This label is added to each array as the 188<sup>th</sup> value. These labels range are either normal or pathological as expressed in Table 1. This type of beat extraction does not filter the signal or make any assumption about the morphology, which has and been shown as an effective method for ECG processing. A further and more detailed overview of these preprocessing steps is provided elsewhere [1].

**Table 1:** The two datasets explored in this project MIT-BIH and PTB Diagnostic

MIT-BIH ECG				PTB Diagnostic ECG			
ECG Class	Label	# samples	% of Total	ECG Class	Label	# samples	% of Total
Normal	0	90589	82.8%	Normal	0	10506	72.2
Fusion of Paced and Normal	1	2779	2.5%	Abnormal	1	4046	27.8
Premature Ventricular Contraction	2	7236	6.6%				
Atrial Premature	3	803	0.7%				
Fusion of Ventricular and Normal	4	8039	7.3%				


### 2.2 Data Augmentation

The number of representative samples for each class shows a large class imbalance (Table 1), where the majority of all data is of normal non-pathological ECG. To address this class imbalance I decided to perform a data augmentation on the minority classes [14]. It is generally suggested that augmenting just the training data increases reliability and generalization of the model [13]. With this in mind, the sample data was first split into 80% training, 20% testing. A *StratifiedShuffleSplit* approach was taken such that the data was split equally amongst the classes. Next, the training data was further split such that 16% went into a separate validation set. The validation and testing sets were then set aside well a very simple data augmentation was performed on the training set for the MIT-BIH dataset (the main dataset explored in the bulk of this project). The majority class (normal ECG) was first reduced down to 10,000 samples and the minority classes were *resampled* such that each class now contained the same amount of observations. Random Gaussian noise was injected to each sample of the resampled data to further increase generalizability. The amplitude of this noise was small (0.03-0.05) to help ensure each label kept its class (it is possible though that this could change the class, this is why a small amount of noise was used). A figurative representation of the data is shown Figure 2 and Table 2 shows the summary of the data after splitting and augmenting.

**Table 2:** The split MIT-BIH dataset showing the number of samples in each class as well the new formed augmented data set which down sampled the majority class and up sampled the minority.

		MIT-BIH ECG Dataset			
	% samples	Training Set 1			Training Set 2
		Non-Augmented Training	Validation Set	Test Set	Augmented Training
		64%	16%	20%	64%
Label number		# samples	# samples	# samples	# samples
Normal	0	57976	14495	18118	10000
Pathological	1	1778	445	556	10000
	2	4631	1158	1447	10000
	3	514	128	161	10000
	4	5145	1286	1608	10000

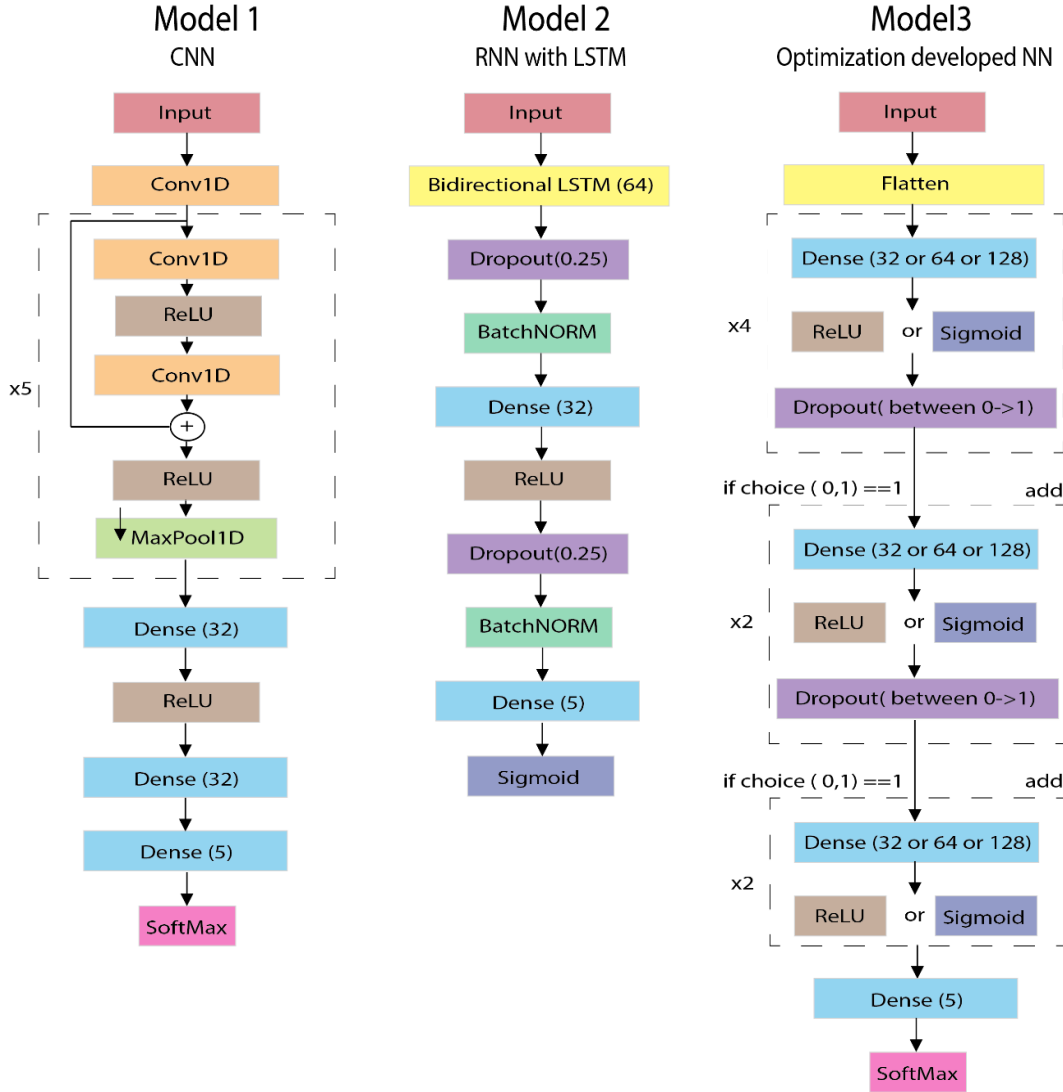
**Figure 2:** Raw data (MIT-BIH) from each category and data with Gaussian noise (augmented data).



### 2.3 Network Architectures

All networks were trained using the following set parameters: batch size = 64, patience = 20, epochs = 500. Model weights were set to be balanced meaning that for the augmented data these weights were just 1 but for the unbalanced data, the weights were balanced following the default approach proposed in Keras using callback functions. Learning rate was programed to have an exponential decay such that it started at 0.001, had 10,000 decay steps at a decay rate of 0.75. The Adam optimizer was used with beta\_1 = 0.9 and beta\_2, 0.999 set. The optimizer optimized for *sparse categorical cross entropy* which is the recommended loss for this type of data [1].

Model 1 is a convolutional neural network with 55,013 parameters based on the architecture proposed elsewhere [1]. Model 2 is a recurrent neural network (RNN) with long short term memory (LSTM) having 54,661 parameters. This model was developed following recommendations for building RNN's and was hand tuned by trial and error during pilot experiments to include dropout [12]. RNN's are well suited for time series data of this type and should work well in classifying the ECG signals [12], [15]. Model 3 is a deep neural network with variable layers that I developed. Since developing neural network architecture is not my expertise, I wanted to use numerical optimization to help develop a simple deep neural network. This process is further explained in section **2.4 Optimization**. The model that trained on the non-augmented data converged on an architecture having 65,669 parameters. The model that trained on the augmented data converged on an architecture having 75,333 parameters (specific models details are provided in the Appendix).



**Figure 3: Model 1:** Convolutional neural network based on [1], **Model 2:** recurrent neural network (RNN) with long short term memory (LSTM), and **Model 3:** an optimization developed neural network with variables parameter and number of layers developed using package *hyperopt* [16].

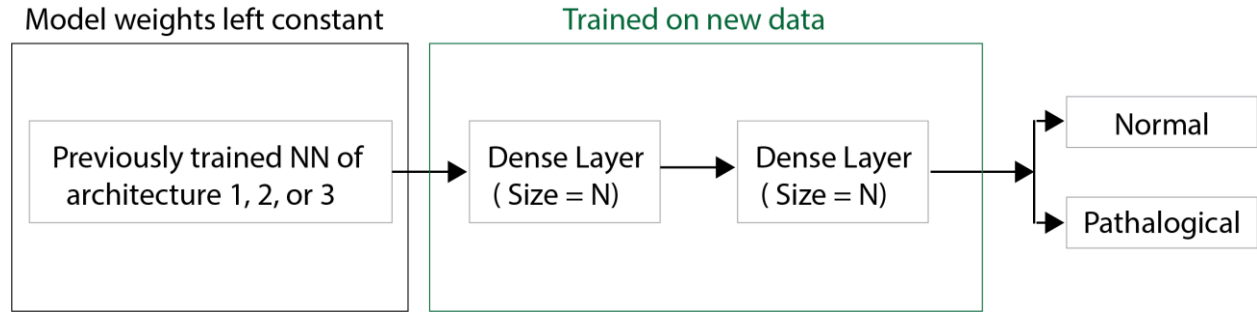
## 2.4 Optimization

An architecture optimization approach was implemented for Model 3 (Figure 3). A numerical optimization was setup to minimize negative test accuracy as depicted by equation 1. The optimization routine first evaluated a possible model architecture based on possible options, then model parameters were generated and optimized by minimizing the sparse *categorical cross entropy*, same as for the other models. The optimization model was improved over 40 epochs of optimization (i.e. 40 different possible architectures were explored and optimized). The best combination of model architecture and model parameters resulting in the lowest objective function (Equation 1) was chosen for further evaluation (resulting model specifics shown in Appendix). This optimization routine was setup up using *hyperopt* [16].

**Equation 1.** 
$$\text{Objective Function} = \min(-(test_{set} \text{ Accuracy}))$$

## 2.5 Transfer Learning

After training and evaluating the models on MIT-BIH dataset, I wanted to see how well the trained model would perform on a novel dataset namely the PTB Diagnostic, which has only two classes: normal and pathological (all pathological signals grouped into one) (Table 1). To do this, I used the best trained model from each architecture and added two dense layers and an output node such that the models now went from the 5 output nodes (5 classes in the MIT-BIH dataset) through two dense layers and outputted either 0 or 1 (normal or pathological) as depicted in Figure 4. This bottlenecking approach was based on methodology proposed in literature [1]. The two added dense layers were then trained on the new data while the other model weights were left untouched.



**Figure 4:** Approach used to add two additional dense layers to previously evaluated models. After some pilot experiments, the added dense layer was set to be a size of  $N = 32$  as is also recommended elsewhere [1].

In addition to this approach, I also evaluated how training deep layers after adding the two dense layers would affect the transfer learning. To evaluate this in a systematic way I also evaluated the effects of training 5 layers deep and training 8 layers deep and present this in my results.

## 2.6 Metrics

The following metrics were utilized in this study. Recall which reflects the accurate positive identification of the ECG label as defined by equation 2.

**Equation 2.**

$$\text{Recall} = \frac{\text{True Positive}}{\text{True Positive} + \text{False Negative}}$$

Precision which is the accurate positive identification—when the model predicted positive was it correct? This is defined by equation 3.

**Equation 3.**

$$\text{Precision} = \frac{\text{True Positive}}{\text{True Positive} + \text{False Positive}}$$

The F1-score is a metric for comparing the performance as a ratio of the mean of the recall and precision. The higher the F1-score the better the model did is the best way to look at it. F1-score is convenient to look at as it can be calculated for each label. In the context of our problem here this is a great metric to use to compare model performance as expressed in equation 4.

**Equation 4.**

$$\text{F1 Score} = \frac{2 * \text{Recall} * \text{Precision}}{\text{Recall} + \text{Precision}}$$

The accuracy is the ratio between correctly predicted outcomes and the sum of all predictions. A relatively easy metric to understand for looking at global performance of the model as depicted by equation 5.

**Equation 5.** 
$$\text{Accuracy} = \frac{\text{True Positive} + \text{True Negative}}{\text{True Positive} + \text{True Negative} + \text{False Positive} + \text{False Negative}}$$

Finally, the error matrix also known as a confusion matrix is a convenient table that shows the percentage of classified or misclassified labels in a matrix (it is the recall measure expressed for every label). Figure 5 provides an example of how this table will be presented in the results of this research. The diagonals of the matrix express the accurate identification of the correct label and the closer this accuracy is to 100% the better the model did. Any numbers shown off the diagonal can be thought of as unfavourable as this expresses misclassification of the model.

Normal	0				
Pathological	1				
	2				
	3				
	4				
		0	1	2	3

**Figure 5:** Example of an error matrix used for evaluating performance of a classification model. High accuracy scores across the diagonal indicates good classification.

## 3.0 Results

### 3.1 Network architecture influences performance

Network architecture has a noticeable effect on the F1-score of each model, as can be seen in Table 3. We can see that Model A has the highest overall F1-score when the data that the network is trained on is augmented (Model A, F1-Score = 0.954) but when the data the network is trained on is not augmented Model B performs best (Model B, F1-Score = 0.993). We can further see that the models all have high accuracies meaning that their ability to perform this classification task is good. The one caveat here is that the testing data is heavily skewed towards normal ECG signals. If the model correctly classifies the normal ECG it can still achieve a high accuracy without correctly classifying pathological ECG.

To better visualize the errors and misclassification occurring in the models, Figure 6 shows the error matrixes. Here we can see that all the models had a difficult time with the pathological label 1(fusion of paced and normal) as indicated by the higher misclassifications (10-35%) of this label. The models performed exceptionally well at classifying the normal ECG signal which was especially noticeable when the data was trained on the non-augmented training set (achieving almost 100% for all three architectures). The ability to well classify the larger normal set also explains why we see such large model accuracies. In general, all the architectures were good at classifying the data regardless of the training dataset however, Model 1 one performed noticeably



better than the other two models. Model 3 had misclassifications errors higher than other models when it came to classifying pathological ECG signals vs. normal.

**Table 3:** F1-scores for the 3 models trained on augmented and non-augmented datasets MIT BIT. Green shading indicates best performance. The closer the F1-score is to 1 the better the model performed at classifying.

		F1-Score			
		Augmented Training	Model 1	Model 2	Model 3
Normal	Normal		0.974	0.960	0.966
	Fusion of Paced and Normal		0.680	0.588	0.679
Pathological	Premature Ventricular Contraction		0.934	0.891	0.847
	Atrial Premature		0.474	0.393	0.484
	Fusion of Ventricular and Normal		0.975	0.972	0.910
Weighted Average			0.960	0.944	0.943
Macro Average			0.807	0.761	0.777

		Non-Augmented Training	Model 1	Model 2	Model 3
Normal	Normal		0.992	0.997	0.989
	Fusion of Paced and Normal		0.843	0.925	0.770
Pathological	Premature Ventricular Contraction		0.961	0.984	0.936
	Atrial Premature		0.799	0.895	0.757
	Fusion of Ventricular and Normal		0.989	0.997	0.982
Weighted Average			0.985	0.993	0.978
Macro Average			0.917	0.960	0.887

**Table 4:** Global accuracy metrics for the 3 models trained on augmented and non-augmented datasets. Green shading indicates best performance. If the model correctly classifies the normal ECG it can still achieve a high accuracy without correctly classifying pathological ECG.

Accuracy	Model 1	Model 2	Model 3
Augmented Training	95.4%	93.2%	97.9%
Non-Augmented Training	98.5%	99.3%	94.0%

Training on Augmented Data																		
# of samples			Model 1 CNN						Model 2 RNN					Model 3 Optimization				
Normal	18118	0	95%	2%	1%	2%	0%	93%	3%	1%	2%	0%	96%	1%	1%	1%	1%	
	556	1	10%	89%	1%	0%	0%	10%	88%	2%	0%	0%	29%	68%	3%	0%	0%	
	1447	2	1%	0%	96%	2%	0%	1%	1%	94%	3%	0%	11%	0%	86%	2%	0%	
	161	3	2%	1%	6%	92%	0%	4%	0%	1%	94%	0%	12%	0%	4%	84%	0%	
	1608	4	2%	0%	1%	0%	97%	1%	0%	0%	0%	99%	4%	0%	2%	0%	94%	


Training on Non-Augmented Data																		
# of samples			Model 1 CNN						Model 2 RNN					Model 3 Optimization				
Normal	18118	0	99%	0%	0%	0%	0%	100%	0%	0%	0%	0%	100%	0%	0%	0%	0%	
	556	1	13%	87%	0%	0%	0%	11%	88%	1%	0%	0%	35%	64%	1%	0%	0%	
	1447	2	3%	1%	94%	1%	0%	1%	0%	98%	1%	0%	6%	0%	91%	3%	0%	
	161	3	12%	2%	6%	80%	0%	7%	0%	6%	88%	0%	20%	0%	3%	76%	0%	
	1608	4	1%	0%	0%	0%	99%	0%	0%	0%	0%	99%	3%	0%	0%	0%	97%	
			0	1	2	3	4	0	1	2	3	4	0	1	2	3	4	

**Figure 6:** Error matrixes for the 3 models explored showing the results for models trained on augmented and non-augmented training sets. Blue shading indicates diagonal squares (100% is perfect classification), red shaded color indicates noticeable misclassifications, grey shading indicates low error (<2%).



### 3.2 Data augmentation decreased performance

Data augmentation had a determinable effect on overall model performance as can be seen by looking at accuracy and loss performance for the three evaluated models shown in Figure 7. The model accuracy was lower and validation loss was higher for models that were trained on the augmented data set resulting in poorer F1-scores (Table 3), lower test accuracy (Table 4) and higher misclassification rates (Figure 6).

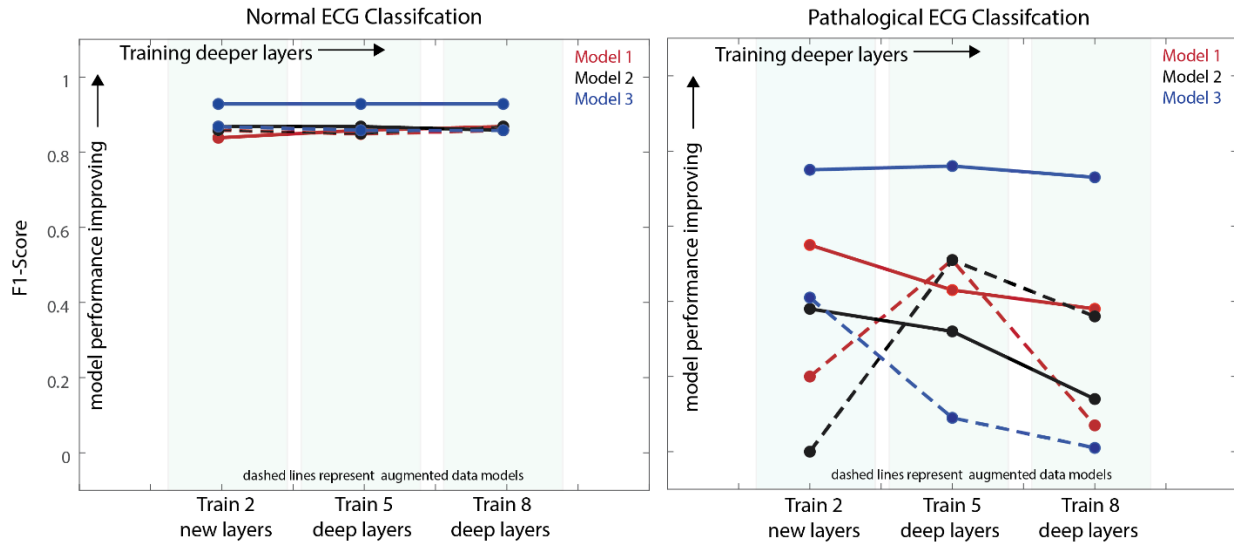


**Figure 7:** Accuracy (top section) and objective function loss (lower section) for validation and training sets using augmented (dashed lines) and non-augmented (solid line) training sets.

The augmented dataset required down sampling the majority class and up sampling the minority class. Perhaps the approach taken using *resample* from the Keras toolbox and adding Gaussian noise was too simple with too many similar cases for the network to train on. It is important to note that there were some improvements on classifying the pathological classes when the models trained on the augmented data (Figure 6) but in general the models performed more poorly. More sophisticated augmentation techniques appear to be beneficial for ECG data classification and should be explored in future work [11].

### 3.3 Transfer learning classification is strong on normal ECG but not on pathological

Figure 8 shows the F1-scores for the transfer learning results comparing: training the 2 added layers, training 5 layers deep, and training up to 8 layers deep. Applying the trained MIT-BIH models on a novel dataset resulted in a strong ability to classify the normal ECG signals but the ability to classify the pathological signal varied widely between models. Classifying the pathological data was particularly poor for the models trained on the augmented dataset. In terms of the F1-score, we can see in Figure 8 that model 3 had the best generalizability to this new dataset, resulting in the highest overall F1-score. It is interesting to note that model 3 had the poorest (although was still good at) classification performance for the MIT-BIH dataset.



**Figure 8:** F1-scores for classification on PTB dataset after transfer learning from MIT\_BIT trained models. PTB has two labels datasets either Normal (right) or Pathological (left). Results for model trained on augmented (dashed lines) and non-augmented (solid lines) training set are shown.

Looking further at the resulting error matrix just for the models trained on the 2 new added layers, we can see that model 3 shows a strong ability to correctly classifying the data, with Model 2 performing poorly (Figure 9). Training more layers did not appear to improve Model 3's performance and in most cases it deteriorated all models ability to classify the pathological data. In general though these results are encouraging and with further tuning improvements could further be made [1], [17].

Training on Augmented Data									
# of samples			Model 1 CNN		Model 2 RNN		Model 3 Opt.		
Normal	2102	0	98%	2%	100%	0%	93%	7%	
Pathological	809	1	12%	88%	100%	0%	69%	31%	
Training on Non-Augmented Data									
			Model 1 CNN		Model 2 RNN		Model 3 Opt.		
Normal	2102	0	82%	18%	94%	6%	92%	8%	
Pathological	809	1	45%	55%	73%	27%	29%	71%	
			0	1	0	1	0	1	

**Figure 9:** Error matrixes for the 3 models explored showing the transfer learning results for models trained on augmented and non-augmented training sets. Blue shading indicates diagonal squares, red shaded color indicates noticeable misclassifications, grey shading indicates low error (<2%). The PTB dataset has two classes normal and pathological. This requires the MIT-BIH trained models to bottle neck down to 2 outputs. Note: these error matrixes are only representative of results from training 2 new added dense layers.

### 3.4 Hyper parameters effect individual model performance

Although all of the results presented in this project used one set of hyper parameters as described in the methods, various other hyper parameter combinations were evaluated during the pilot testing of developing this systematic experiment. It warrants to present the effects of hyper parameter tuning to increase the transparency and further possibility of improving the presented results. Changes in parameters such as batch size, augmentation sample set or patience each resulted in

differences in the final outcomes, as presented by the error matrix of an arbitrary model architecture with changes made to these parameters in Figure 10. Although further tuning and exploring the best combinations of these parameters is possible, there were limitations in the hardware the network was running on and the time constraints possible to evaluate long duration computations (Anaconda using a single-core 3.6 GHz Dell OptiPlex 7020 computer with 16 Gb of RAM). I am acknowledging that it is possible that the results presented here could be further improved upon, but I feel that the conclusions and end outcomes would likely not vary drastically from what was presented here. Nevertheless, further work could lead to improvements in performance and should be explored.

Evaluation of hyper-parameter tuning on outcomes																			
		Batch size = 256						Augmentation = 20,000						Patience = 10					
Normal	18118	0	82%	14%	1%	2%	1%	98%	2%	1%	0%	0%	93%	5%	1%	0%	0%		
	556	1	5%	94%	1%	0%	0%	9%	90%	1%	0%	0%	3%	96%	1%	0%	0%		
Pathological	1447	2	2%	3%	91%	4%	1%	2%	1%	96%	1%	0%	1%	1%	97%	1%	0%		
	161	3	0%	4%	2%	94%	0%	9%	1%	7%	83%	0%	5%	2%	11%	82%	0%		
	1608	4	0%	0%	0%	0%	99%	1%	0%	0%	0%	99%	0%	0%	0%	0%	99%		
			0	1	2	3	4	0	1	2	3	4	0	1	2	3	4		

**Figure 10:** Error matrixes for 1 example model using variations in hyper parameters. Results show sensitivity to user selected hyper parameters, further work is warranted.

### 3.5 Comparison of results to literature

The results presented in this project were compared to results from other sources. Specially, the accuracy and F1-score was compared as shown in Table 5. The performance metrics achieved in this work are strong and compare well against what is found in recent literature. Although the classification ability of the proposed work is certainly good, caution must be taken when making direct comparisons of this sort. The test set used in this project was unbalanced, small, and only had several pathological signals (for example [3][7] evaluated 12 classes on over 50,000 patients). Making comparison to other models which may achieve lower F1-score and lower accuracy for this reason is deceptive and must be taken in to account. For this specific problem misclassification of pathological ECG has larger consequences then misclassifying a healthy normal ECG. Although this was not specifically taken into account in this project it is important to consider when evaluating the models ability to classify ECG signals.

**Table 4:** Comparing the accuracy and F1-scores obtained in this project to that of literature

Authors	Citation	accuracy %	Architecture	Authors	Citation	F1-score	F1-score (cardiologist)
2019 Mousavi et al.	[13]	99.92*	RNN+CNN	2020 Steenkiste et al.	[17]	0.80-0.98	
2017 Achaya et al.	[11]	97.40	CNN	2019 Hannun et al.	[3][7]	0.57-0.94**	0.53-0.91**
2016 Kiranyaz et al.	[6]	99.10*	CNN 2D	2019 Alfaras et al.	[10]	0.78-0.98	
2010 Ye et al.	[8]	99.91*	Support Vector Machine				
Presented here		93.2-99.3		Presented here		0.39-0.99	

\* more ECG classes explored then in this dataset and large sample size with augmentation

\*\* Over 12 classes and dataset composed of over 50,000+ patients

## 4.0 Conclusions

In summary, the work presented here successfully met the objectives of the project. First, I was able to implement 3 neural network architectures for classifying ECG signals. All the proposed architectures had strong performance metrics and were able to classify the data to a high degree of accuracy. Second, I augmented the training set data and evaluated the performance effects. I found that the simple data augmentation approach led to an overall decrease in prediction performance. Lastly, I evaluated how well the models would perform when tested on a new data set. I found that the classification of normal ECG was strong but the classifying the pathological data proved to be more difficult. Further work is warranted in optimizing hyper parameters which would lead to improvements in the outcomes.

## References

- [1] M. Kachuee, S. Fazeli, and M. Sarrafzadeh, "ECG Heartbeat Classification: A Deep Transferable Representation," *2018 IEEE Int. Conf. Healthc. Inform. ICHI*, pp. 443–444, Jun. 2018, doi: 10.1109/ICHI.2018.00092.
- [2] S. Kiranyaz, T. Ince, and M. Gabbouj, "Real-Time Patient-Specific ECG Classification by 1-D Convolutional Neural Networks," *IEEE Trans. Biomed. Eng.*, vol. 63, no. 3, pp. 664–675, Mar. 2016, doi: 10.1109/TBME.2015.2468589.
- [3] A. Y. Hannun *et al.*, "Cardiologist-level arrhythmia detection and classification in ambulatory electrocardiograms using a deep neural network," *Nat. Med.*, vol. 25, no. 1, pp. 65–69, Jan. 2019, doi: 10.1038/s41591-018-0268-3.
- [4] S. M. Jadhav, S. L. Nalbalwar, and A. A. Ghatol, "ECG arrhythmia classification using modular neural network model," in *2010 IEEE EMBS Conference on Biomedical Engineering and Sciences (IECBES)*, Nov. 2010, pp. 62–66, doi: 10.1109/IECBES.2010.5742200.
- [5] J. Schl pfer and H. J. Wellens, "Computer-Interpreted Electrocardiograms," *J. Am. Coll. Cardiol.*, vol. 70, no. 9, pp. 1183–1192, Aug. 2017, doi: 10.1016/j.jacc.2017.07.723.
- [6] S. Kiranyaz, T. Ince, R. Hamila, and M. Gabbouj, "Convolutional Neural Networks for patient-specific ECG classification," in *2015 37th Annual International Conference of the IEEE Engineering in Medicine and Biology Society (EMBC)*, Milan, Aug. 2015, pp. 2608–2611, doi: 10.1109/EMBC.2015.7318926.
- [7] P. Rajpurkar, A. Y. Hannun, M. Haghp n hi, C. Bourn, and A. Y. Ng, "Cardiologist-Level Arrhythmia Detection with Convolutional Neural Networks," *ArXiv170701836 Cs*, Jul. 2017, Accessed: Mar. 18, 2020. [Online]. Available: <http://arxiv.org/abs/1707.01836>.
- [8] Can Ye, M. T. Coimbra, and B. V. K. Vijaya Kumar, "Arrhythmia detection and classification using morphological and dynamic features of ECG signals," in *2010 Annual International Conference of the IEEE Engineering in Medicine and Biology*, Buenos Aires, Aug. 2010, pp. 1918–1921, doi: 10.1109/IEMBS.2010.5627645.
- [9] T. Li and M. Zhou, "ECG Classification Using Wavelet Packet Entropy and Random Forests," *Entropy*, vol. 18, no. 8, p. 285, Aug. 2016, doi: 10.3390/e18080285.
- [10] M. Alfaras, M. C. Soriano, and S. Ort n, "A Fast Machine Learning Model for ECG-Based Heartbeat Classification and Arrhythmia Detection," *Front. Phys.*, vol. 7, p. 103, Jul. 2019, doi: 10.3389/fphy.2019.00103.
- [11] U. R. Acharya *et al.*, "A deep convolutional neural network model to classify heartbeats," *Comput. Biol. Med.*, vol. 89, pp. 389–396, Oct. 2017, doi: 10.1016/j.combiomed.2017.08.022.
- [12] S. Saadatnejad, M. Oveisi, and M. Hashemi, "LSTM-Based ECG Classification for Continuous Monitoring on Personal Wearable Devices," *IEEE J. Biomed. Health Inform.*, vol. 24, no. 2, pp. 515–523, Feb. 2020, doi: 10.1109/JBHI.2019.2911367.

- [13] S. Mousavi, F. Afghah, and U. R. Acharya, "Inter- and intra- patient ECG heartbeat classification for arrhythmia detection: a sequence to sequence deep learning approach," *ArXiv181207421 Phys. Q-Bio*, Mar. 2019, Accessed: Mar. 18, 2020. [Online]. Available: <http://arxiv.org/abs/1812.07421>.
- [14] "Resampling strategies for imbalanced datasets." <https://kaggle.com/rafjaa/resampling-strategies-for-imbalanced-datasets> (accessed Apr. 12, 2020).
- [15] M. Zihlmann, D. Perekrestenko, and M. Tschannen, "Convolutional Recurrent Neural Networks for Electrocardiogram Classification," in *Computing in Cardiology*, Sep. 2017, doi: 10.22489/CinC.2017.070-060.
- [16] J. Bergstra, D. Yamins, and D. D. Cox, "Making a Science of Model Search: Hyperparameter Optimization in Hundreds of Dimensions for Vision Architectures," p. 9.
- [17] G. Van Steenkiste, G. van Loon, and G. Crevecoeur, "Transfer Learning in ECG Classification from Human to Horse Using a Novel Parallel Neural Network Architecture," *Sci. Rep.*, vol. 10, no. 1, p. 186, Dec. 2020, doi: 10.1038/s41598-019-57025-2.

# 4.0 Appendix

## Model 3: Trained on non-augmented data set

Options resulting from optimization see Figure 3 for model

characteristics  
{'Activation': 0,  
'Activation\_1': 0,  
'Activation\_2': 0,  
'Activation\_3': 0,  
'Activation\_4': 0,  
'Activation\_5': 1,  
'Activation\_6': 1,  
'Activation\_7': 0,  
'Activation\_8': 0,  
'Dense': 2,  
'Dense\_1': 2,  
'Dense\_2': 2,  
'Dense\_3': 1,  
'Dense\_4': 1,  
'Dense\_5': 1,  
'Dense\_6': 1,  
'Dense\_7': 0,  
'Dense\_8': 1,  
'Dense\_9': 3,  
'Dropout': 0.3265904623792442,  
'Dropout\_1': 0.3431590875721917,  
'Dropout\_2': 0.25572580783477367,  
'Dropout\_3': 0.6975060560039738,  
'Dropout\_4': 0.3423233980953824,  
'Dropout\_5': 0.32317538374749766,  
'a': 0,  
'if': 0,  
'if\_1': 0,  
'if\_2': 0,  
'if\_3': 0}

Model: "sequential_103"		
Layer (type)	Output Shape	Param #
=====		
flatten_107 (Flatten)	(None, 187)	0
dense_677 (Dense)	(None, 128)	24064
activation_636 (Activation)	(None, 128)	0
dropout_456 (Dropout)	(None, 128)	0
dense_678 (Dense)	(None, 128)	16512
activation_637 (Activation)	(None, 128)	0
dropout_457 (Dropout)	(None, 128)	0
dense_679 (Dense)	(None, 128)	16512
activation_638 (Activation)	(None, 128)	0
dropout_458 (Dropout)	(None, 128)	0
dense_680 (Dense)	(None, 64)	8256
activation_639 (Activation)	(None, 64)	0
dropout_459 (Dropout)	(None, 64)	0
dense_681 (Dense)	(None, 5)	325
=====		
Total params: 65,669		

## Model 3: Trained on augmented data set

Options resulting from optimization see Figure 3

for model characteristics	Layer (type)	Output Shape	Param #
{'Activation': 0,	=====		
'Activation_1': 1,	flatten_135 (Flatten)	(None, 187)	0
'Activation_2': 0,			
'Activation_3': 0,	dense_917 (Dense)	(None, 128)	24064
'Activation_4': 0,			
'Activation_5': 0,	activation_815 (Activation)	(None, 128)	0
'Activation_6': 0,			
'Activation_7': 0,	dropout_582 (Dropout)	(None, 128)	0
'Activation_8': 1,			
'Dense': 2,	dense_918 (Dense)	(None, 32)	4128
'Dense_1': 0,			
'Dense_2': 2,	activation_816 (Activation)	(None, 32)	0
'Dense_3': 1,			
'Dense_4': 2,	dropout_583 (Dropout)	(None, 32)	0
'Dense_5': 1,			
'Dense_6': 0,	dense_919 (Dense)	(None, 128)	4224
'Dense_7': 0,			
'Dense_8': 0,	activation_817 (Activation)	(None, 128)	0
'Dense_9': 1,			
'Dropout': 0.4959326042904314,	dropout_584 (Dropout)	(None, 128)	0
'Dropout_1': 0.5162300963476361,			
'Dropout_2': 0.4169915806152422,	dense_920 (Dense)	(None, 64)	8256
'Dropout_3': 0.25301754583882874,			
'Dropout_4': 0.2515810749300903,	activation_818 (Activation)	(None, 64)	0
'Dropout_5': 0.516131118961639,			
'a': 1,	dropout_585 (Dropout)	(None, 64)	0
'if': 0,			
'if_1': 0,	dense_921 (Dense)	(None, 32)	2080
'if_2': 1,			
'if_3': 0}	activation_819 (Activation)	(None, 32)	0
	dense_922 (Dense)	(None, 5)	165
	=====		
	Total params: 75,333		



```

# Written by Pawel Kudzia

import pandas as pd
import numpy as np
import tensorflow as tf
import matplotlib.pyplot as plt

from sklearn.utils import resample
from sklearn.utils import class_weight

from sklearn.model_selection import (TimeSeriesSplit, KFold, ShuffleSplit,
                                     StratifiedKFold, GroupShuffleSplit,
                                     GroupKFold, StratifiedShuffleSplit)

from sklearn.metrics import classification_report, roc_auc_score, accuracy_score, precision_score, recall_score

from tensorflow.keras.layers import Input, Conv1D, MaxPool1D, Activation, Add, Dense, Flatten
from tensorflow.keras.models import Model, load_model
from tensorflow.keras.optimizers import Adam
from tensorflow.keras.utils import plot_model
from tensorflow.keras.utils import to_categorical
from tensorflow.keras.metrics import Precision
from tensorflow.keras.layers import Input, Dropout, LSTM, Bidirectional, BatchNormalization, Dense, Activation
from tensorflow.keras.models import Sequential, Model

from tensorflow.keras.callbacks import EarlyStopping, LearningRateScheduler, ModelCheckpoint

from sklearn.metrics import roc_curve, auc, roc_auc_score

from sklearn.metrics import confusion_matrix

# CUSTOM Functions imported from folder
from plot_Conf_Matrix import plot_Conf_Matrix
from make_model_1 import make_model_1 # This is the biorix Architecture Model 1 in submitted report
from make_model_2_RNN import make_model_2_RNN # this is a simple RNN model

from hyperopt import Trials, STATUS_OK, tpe
from hyperas import optim
from hyperas.distributions import choice, uniform

import warnings
warnings.filterwarnings('ignore')
warnings.simplefilter(action='ignore', category=FutureWarning)

# %% Modelling Parameters and Variables CHANGE PARAMETERS HERE AS NEEDED
saveData = 0

epochs = 1
batch = 64
patience = 10
verbose = 1
ModelSaveName = 'Model1CNN_UsingNotAugmentedData.h5' # Change this to what you want to save the model as
Model2_RNN = 0
Model1_CNN = 1
Model3_OptimizationModel = 0

# %% Class labels
classesMIT={0:"Normal",
            1:"Artial Premature",
            2:"Premature Ventricular Contraction",
            3:"Fusion of Centricular and Normal",
            4:"Fusion of paced and Normal"}

classePTB={0:"Normal", 1:"Abnormal (MI)"}

# %% Let define all our smaller nested functions here

def plot_learning(history):
    plt.subplot(211)
    plt.plot(history.history['accuracy'])
    plt.legend(['accuracy'])
    plt.subplot(212)
    plt.plot(history.history['loss'])
    plt.plot(history.history['val_loss'], label = "val_loss")
    plt.legend(['loss', "val_loss"])
    plt.show()

def outputMetrics(X_test, Y_Test, ModelX, batch) :
    if 1:
        X_test = X_test.reshape(len(X_test), X_test.shape[1],1)
    else:
        X_test = X_test

    Y_Test=to_categorical(Y_Test)

    PredictionModel=ModelX.predict(X_test, batch_size=batch, verbose=1)

    TrueValue=np.argmax(Y_Test, axis=1)
    Prediction=np.argmax(PredictionModel,axis=1)

    print(classification_report(TrueValue, Prediction))

    return (Prediction,TrueValue)

def roc_auc_score_multiclass(actual_class, pred_class, average = "macro"):
    #creating a set of all the unique classes using the actual class list
    unique_class = set(actual_class)
    roc_auc_dict = {}

```

```

102 roc_auc_dict = {}
103 for per_class in unique_class:
104     #creating a list of all the classes except the current class
105     other_class = [x for x in unique_class if x != per_class]
106
107     #marking the current class as 1 and all other classes as 0
108     new_actual_class = [0 if x in other_class else 1 for x in actual_class]
109     new_pred_class = [0 if x in other_class else 1 for x in pred_class]
110
111     #using the sklearn metrics method to calculate the roc_auc_score
112     roc_auc = roc_auc_score(new_actual_class, new_pred_class, average = average)
113     roc_auc_dict[per_class] = roc_auc
114
115     return roc_auc_dict
116
117 def classification_report_csv(report):
118     report_data = []
119     lines = report.split("\n")
120     for line in lines[2:-3]:
121         row = {}
122         row_data = ''.join(line.split())
123         row_data = row_data.split(' ')
124         row['class'] = row_data[0]
125         row['precision'] = float(row_data[1])
126         row['recall'] = float(row_data[2])
127         row['f1_score'] = float(row_data[3])
128         row['support'] = float(row_data[4])
129         report_data.append(row)
130
131     dataframe = pd.DataFrame.from_dict(report_data)
132     dataframe.to_csv('classification_report.csv', index = False)
133
134     return()
135
136 def NetworkModel(ModelInput, TrainingDataX, TrainingDataY, ValidationX, ValidationY, TestSetX, TestSetY, epochs, batch, patience, verbose, ModelSaveName, ComputeWeight, RNN, CNN, Optimization):
137
138     # he learning hyper paramters are just set to what is recommended adjusting these paramters will not be the goal of this project.
139     # these specific paramaters were taken from the paper "ECG Heartbeat Classification: A Deep Transferable Representation"
140     lr_schedule = tf.keras.optimizers.schedules.ExponentialDecay(0.001, decay_steps=10000, decay_rate=0.75)
141     adam = Adam(learning_rate=lr_schedule, beta_1=0.9, beta_2=0.999, amsgrad=False)
142
143     if ComputeWeight:
144         class_weights = class_weight.compute_class_weight('balanced',
145                                                         np.unique(TrainingDataY),
146                                                         TrainingDataY)
147     else:
148         class_weights = []
149
150     # We will use the sparse_C_C and the optimization function and look at accuracy for now
151     ModelInput.compile(optimizer=adam, loss='sparse_categorical_crossentropy', metrics=['accuracy'])
152
153     # Define some early stopping criteria here
154     callb1 = EarlyStopping(monitor='val_loss', mode='min', verbose=verbose, patience=patience) # lets monitor the loss in the opt function
155     callb2 = ModelCheckpoint(filepath=ModelSaveName, monitor='val_loss', save_best_only=True) # Lets save it everytime it improves as well
156
157     if RNN: # fit setup for RNN network
158
159         history = ModelInput.fit(TrainingDataX,
160                                 TrainingDataY,
161                                 epochs=epochs,
162                                 batch_size=batch,
163                                 validation_data=(ValidationX, ValidationY),
164                                 verbose=verbose,
165                                 class_weight = class_weights,
166                                 callbacks=[callb1, callb2])
167
168     elif CNN: # fit setup for CNN network
169
170         history = ModelInput.fit(np.expand_dims(TrainingDataX, axis=2),
171                                 TrainingDataY,
172                                 epochs=epochs,
173                                 batch_size=batch,
174                                 validation_data=(np.expand_dims(ValidationX, axis=2), ValidationY),
175                                 verbose=verbose,
176                                 class_weight = class_weights,
177                                 callbacks=[callb1, callb2])
178
179     plot_learning(history)
180     ModelInput.save(ModelSaveName)
181     print("The model has been saved into your directory")
182
183
184     return(ModelInput, history)
185
186 def optimization_model(TrainingDataX, TrainingDataY, ValidationX, ValidationY, X_Test_MIT, Y_Test_MIT):
187
188     ModelForTraining = Sequential()
189     ModelForTraining.add(Flatten(input_shape=(187, 1)))
190
191     ModelForTraining.add(Dense([choice([32, 64, 128])], input_shape=(187, 1)))
192     ModelForTraining.add(Activation([choice(['relu', 'sigmoid'])]))
193     ModelForTraining.add(Dropout([uniform(0.25, 0.75)]))
194
195     ModelForTraining.add(Dense([choice([32, 64, 128])], input_shape=(187, 1)))
196
197     ModelForTraining.add(Activation([choice(['relu', 'sigmoid'])]))
198     ModelForTraining.add(Dropout([uniform(0.25, 0.75)]))
199
200     ModelForTraining.add(Dense([choice([32, 64, 128])], input_shape=(187, 1)))
201     ModelForTraining.add(Activation([choice(['relu', 'sigmoid'])]))
202     ModelForTraining.add(Dropout([uniform(0.25, 0.75)]))
203
204     ModelForTraining.add(Dense([choice([32, 64, 128])], input_shape=(187, 1)))
205     ModelForTraining.add(Activation([choice(['relu', 'sigmoid'])]))
206     ModelForTraining.add(Dropout([uniform(0.25, 0.75)]))

```

```

207
208
209 if {{{choice(['two', 'three'])}}} == 'three':
210     ModelForTraining.add(Dense({{choice([32,64,128])}}))
211     ModelForTraining.add(Activation({{choice(['relu', 'sigmoid'])}}))
212     ModelForTraining.add(Dropout({{uniform(0.25, 0.75)}}))
213
214
215 if {{{choice(['two', 'three'])}}} == 'three':
216     ModelForTraining.add(Dense({{choice([32,64,128])}}))
217     ModelForTraining.add(Activation({{choice(['relu', 'sigmoid'])}}))
218     ModelForTraining.add(Dropout({{uniform(0.25, 0.75)}}))
219
220
221 if {{{choice(['two', 'three'])}}} == 'three':
222     ModelForTraining.add(Dense({{choice([32,64,128])}}))
223     ModelForTraining.add(Activation({{choice(['relu', 'sigmoid'])}}))
224
225
226 if {{{choice(['two', 'three'])}}} == 'three':
227     ModelForTraining.add(Dense({{choice([32,64,128])}}))
228     ModelForTraining.add(Activation({{choice(['relu', 'sigmoid'])}}))
229
230
231 ModelForTraining.add(Dense(5,activation='softmax'))
232
233 # in pilot testing i was experimenting with optimizing Model 1 to understand if the authors in the paper used the best paramters i.e if i could achieve better results
234 # although this was a good effort the processing was simply too slow on my hardware to make this a
235
236 # input_shape = (187, 1)
237 # l = Input(input_shape)
238 # C = Conv1D(filters=32, kernel_size=5 ,strides=1)(l)
239
240 # a = {{choice(['relu', 'sigmoid','softmax'])}}
241
242 # C11 = Conv1D(filters=32, kernel_size=5,strides=1,padding='same')(C)
243 # R11 = Activation(activation=a)(C11)
244 # C12 = Conv1D(filters=32, kernel_size=5,strides=1,padding='same')(R11)
245 # A11 = Add()(C12, C)
246 # R12 = Activation(activation=a)(A11)
247 # M11 = MaxPool1D(pool_size=5, strides=2)(R12)
248
249 # C21 = Conv1D(filters=32, kernel_size=5,strides=1,padding='same')(M11)
250 # R21 = Activation(activation= a)(C21)
251 # C22 = Conv1D(filters=32, kernel_size=5,strides=1,padding='same')(R21)
252 # A21 = Add()(C22, M11)
253 # R22 = Activation(activation=a)(A21)
254 # M21 = MaxPool1D(pool_size=5, strides=2)(R22)
255
256 # C31 = Conv1D(filters=32, kernel_size=5,strides=1,padding='same')(M21)
257 # R31 = Activation(activation=a)(C31)
258 # C32 = Conv1D(filters=32, kernel_size=5,strides=1,padding='same')(R31)
259 # A31 = Add()(C32, M21)
260 # R32 = Activation(activation=a)(A31)
261 # M31 = MaxPool1D(pool_size=5, strides=2)(R32)
262
263 # C41 = Conv1D(filters=32, kernel_size=5,strides=1,padding='same')(M31)
264 # R41 = Activation(activation=a)(C41)
265 # C42 = Conv1D(filters=32, kernel_size=5,strides=1,padding='same')(R41)
266 # A41 = Add()(C42, M31)
267 # R42 = Activation(activation=a)(A41)
268 # M41 = MaxPool1D(pool_size=5, strides=2)(R42)
269
270 # C51 = Conv1D(filters=32, kernel_size=5,strides=1,padding='same')(M41)
271 # R51 = Activation(activation=a)(C51)
272 # C52 = Conv1D(filters=32, kernel_size=5,strides=1,padding='same')(R51)
273 # A51 = Add()(M41,C52)
274 # R52 = Activation(activation=a)(A51)
275 # M51 = MaxPool1D(pool_size=5, strides=2)(R52)
276
277 # F1 = Flatten()(M51)
278
279 # D1 = Dense({{choice([32,64,128,256])}})(F1)
280 # R6 = Activation(activation={{choice(['relu', 'sigmoid'])}})(D1)
281 # D2 = Dense({{choice([32,64,128,256])}})(R6)
282 # D3 = Dense(5)(D2)
283 # O = Activation(activation='softmax')(D3)
284
285 # ModelForTraining = Model(inputs=l, outputs=O)
286
287 lr_schedule = tf.keras.optimizers.schedules.ExponentialDecay(0.001, decay_steps=10000, decay_rate=0.75)
288 adam = Adam(learning_rate=lr_schedule, beta_1=0.9, beta_2=0.999, amsgrad=False)
289
290 class_weights = class_weight.compute_class_weight('balanced',
291                                                    np.unique(TrainingDataY),
292                                                    TrainingDataY)
293
294 # class_weights = {0: 1., 1: 1.5, 2: 1.5, 3:1.5, 4:1.5} Experimenting here with different weights
295
296 # We will use the sparse_C_C and the optimization function and look at accruacy for now
297 ModelForTraining.compile(optimizer=adam, loss='sparse_categorical_crossentropy', metrics=['accuracy'])
298
299 # Define some early stopping criteria here (this was not used for optimization as it resulted in poor results)
300 # callb1= EarlyStopping(monitor='val_loss', mode='min', verbose=1,patience=10) # lets monitor the loss in the opt function
301 # callb2 = ModelCheckpoint(filepath=ModelSaveName, monitor='val_loss', save_best_only=True) # Lets save it everytime it imprves as well
302
303 history = ModelForTraining.fit(np.expand_dims(TrainingDataX, axis=2),
304                               TrainingDataY,
305                               epochs=30,
306                               batch_size=64,
307                               validation_data=(np.expand_dims(ValidationX, axis=2), ValidationY),
308                               verbose=0,
309                               class_weight=class_weights)
310 # callbacks=[callb1])
311
312 test_loss, test_acc = ModelForTraining.evaluate(np.expand_dims(X_Test_MIT,axis=2), Y_Test_MIT, batch_size=64)

```

```

313 #get the highest validation accuracy of the training epochs
314 # validation_acc = np.amax(history.history['val_accuracy']) # Not used
315
316 print('Best validation acc of epoch:', test_acc)
317
318 return {'loss': -test_acc, 'status': STATUS_OK, 'model': ModelForTraining}
319
320 def CalculateModelMetrics(ModelOutput,History,TestSetX,TestSetY,batch, Naming):
321
322     # function to make excel sheets and plot some of the data for further inspection
323
324     # results = ModelOutput.evaluate(np.expand_dims(TestSetX, axis=2), TestSetY, batch_size=batch)
325     # print(f"The accuracy on the testing set is {np.round(results[1]*100, 1)}%")
326
327     [Prediction,TrueValue,] =outputMetrics(TestSetX,TestSetY,ModelOutput,batch)
328     clsf_reportClassification = pd.DataFrame(classification_report(y_true =TrueValue, y_pred = Prediction, output_dict=True)).transpose()
329
330     confusionM = confusion_matrix(TrueValue, Prediction)
331     clsf_report = pd.DataFrame(confusionM)
332
333     plt.figure()
334     ReportName = "%s" %(Naming)
335
336     plot_Conf_Matrix(confusionM,classes=['N', 'S', 'V', 'F', 'Q'],normalize=True,title=ReportName)
337     #plot_Conf_Matrix(confusionM,classes=['N', 'S'],normalize=True,title=ReportName)
338     plt.show()
339
340     print("\nLogistic Regression")
341     # assuming your already have a list of actual_class and predicted_class from the logistic regression classifier
342     lr_roc_auc_multiclass = roc_auc_score_multiclass(TrueValue, Prediction)
343     print(lr_roc_auc_multiclass)
344
345     hist_df = pd.DataFrame(History.history)
346
347     hist_csv_file = ' History_%s' %(Naming)
348     with open(hist_csv_file, mode='w') as f:
349         hist_df.to_csv(f)
350
351     jsonName = Naming.replace('.csv', '')
352     hist_json_file = ' History_%s.json' %(jsonName)
353     with open(hist_json_file, mode='w') as f:
354         hist_df.to_json(f)
355
356     ReportName = 'ConfusionMatrix%s' %(Naming)
357     clsf_report.to_csv(ReportName, index= True)
358
359     ReportName = 'ClassificationReport%s' %(Naming)
360     clsf_reportClassification.to_csv(ReportName, index= True)
361     return()
362
363
364 def data_curation_MIT():
365     # this function imports the data, performs the augmentation etc.
366
367     SampleCounts= 10000
368     noiseLevel = 0.03
369
370     csv_MIT_train = pd.read_csv(r'E:\1_PAWEL\ENSC_project\heart_beatdata\mitbih_train.csv',header=None)
371     csv_MIT_test= pd.read_csv(r'E:\1_PAWEL\ENSC_project\heart_beatdata\mitbih_test.csv',header=None)
372     DataSet_MIT= pd.concat([csv_MIT_train, csv_MIT_test], axis=0) # put it all into one structure
373     Values_MIT = DataSet_MIT.values
374
375     X_Data_MIT = Values_MIT[:,:-1] # remove the last column which is the calssification
376     Y_Data_MIT = Values_MIT[:,1] # keep the last column which will now be used for the Class Categoization
377
378     StratifyData = StratifiedShuffleSplit(n_splits=5, test_size=0.2, random_state=100)
379     # what we are trying to accomplish here is to split up the data such that the number of samples in each class is somewhat evenly distributed
380     # Since we already have a large class imbalance , this might be an issue if it is not addressed from the start. Lets split it up such that the
381     # test_size get 20% of the data (but catigorical samples evenly distributed)
382
383     StratifyData.get_n_splits(X_Data_MIT, Y_Data_MIT)
384
385     print(StratifyData)
386
387     for train_index, test_index in StratifyData.split(X_Data_MIT, Y_Data_MIT):
388         X_temp, X_Test_MIT = X_Data_MIT[train_index], X_Data_MIT[test_index]
389         y_temp, Y_Test_MIT = Y_Data_MIT[train_index], Y_Data_MIT[test_index]
390
391     # Now lets further split up our training data such that 20% of it goes to the validation sample
392     #[X_Train_MIT, X_Valid_MIT, Y_Train_MIT, Y_Valid_MIT] = train_test_split(X_temp, y_temp, test_size=0.2)
393
394     StratifyData = StratifiedShuffleSplit(n_splits=5, test_size=0.20, random_state=100)
395     # what we are trying to accomplish here is to split up the data such that the number of samples in each class is somewhat evenly distributed
396     # Since we already have a large class imbalance , this might be an issue if it is not addressed from the start. Lets split it up such that the
397     # test_size get 20% of the data (but catigorical samples evenly distributed)
398     StratifyData.get_n_splits(X_temp, y_temp)
399
400     for train_index, test_index in StratifyData.split(X_temp, y_temp):
401         X_Train_MIT, X_Valid_MIT = X_temp[train_index], X_temp[test_index]
402         Y_Train_MIT, Y_Valid_MIT = y_temp[train_index], y_temp[test_index]
403
404     # TrainingDataX = X_Train_MIT # TRAINING DATA
405     # TrainingDataY = Y_Train_MIT# TRAINING DATA Labels
406
407     ValidationX =X_Valid_MIT
408     ValidationY =Y_Valid_MIT
409
410     temp_X_Train_MIT = pd.DataFrame(X_Train_MIT)
411     temp_Y_Train_MIT = pd.DataFrame(Y_Train_MIT)
412
413     tempMIT_TrainingDataSet= pd.concat([temp_X_Train_MIT,temp_Y_Train_MIT], axis=1,ignore_index=True) # put it all into one structure, ignore the index and autogenerate a new one
414
415     print(tempMIT_TrainingDataSet[187].astype(int).value_counts()) # Lets see how many there are in each category. # so we have ~50k in the majority and >5k in the minoritys
416     # at this point i think it makes sense to slightly reduce the sample count on the majority and then upsample the minority so ...
417
418

```

```

410 temp1=tempMIT_TrainingDataSet[tempMIT_TrainingDataSet[187]==1]
411 temp2=tempMIT_TrainingDataSet[tempMIT_TrainingDataSet[187]==2]
412 temp3=tempMIT_TrainingDataSet[tempMIT_TrainingDataSet[187]==3]
413 temp4=tempMIT_TrainingDataSet[tempMIT_TrainingDataSet[187]==4]
414
415 temp0=(tempMIT_TrainingDataSet[tempMIT_TrainingDataSet[187]==0]).sample(n=SampleCounts,random_state=1) # We can downsample this
416 temp1_upsample=resample(temp1,replace=True,n_samples=SampleCounts,random_state=2) # and we can upsample all of the minority classes
417 temp2_upsample=resample(temp2,replace=True,n_samples=SampleCounts,random_state=3) # one note here is the random_state=var ; this will let us reproduce the results everytime we run this
418 temp3_upsample=resample(temp3,replace=True,n_samples=SampleCounts,random_state=4)
419 temp4_upsample=resample(temp4,replace=True,n_samples=SampleCounts,random_state=5)
420
421 # note that this resample approach will just randomly duplicate observations from the minority class. we might want to add some noise to
422 # this data to avoid the neural network just learning the samples. However, we need to be careful here as to not physically change the sample and its
423 # classification. Therefore we will just add a bit of gaussian noise
424
425 MIT_Training_Resampled=pd.concat([temp0,temp1_upsample,temp2_upsample,temp3_upsample,temp4_upsample])
426 print(MIT_Training_Resampled[187].astype(int).value_counts())
427
428 MIT_Training_Resampled = MIT_Training_Resampled.values
429 X_Data_MIT_resampled_Training = MIT_Training_Resampled[:,:-1] # remove the last column which is the classification
430 Y_Data_MIT_resampled_Training = MIT_Training_Resampled[:, -1] # keep the last column which will now be used for the Class Categorization
431
432 X_Data_MIT_resampled_Noise_Training = np.empty(X_Data_MIT_resampled_Training.shape)
433
434 for i in range(0,len(X_Data_MIT_resampled_Training)):
435     noise1=np.random.normal(0,noiseLevel,187)
436     noise2=np.random.normal(0,noiseLevel+0.01,187)+noise1
437     noise=np.random.normal(0,noiseLevel+0.01,187)+noise2
438     X_Data_MIT_resampled_Noise_Training[i] = X_Data_MIT_resampled_Training[i]+noise # lets loop through each sample and add some random noise
439
440
441 TrainingDataX = X_Data_MIT_resampled_Training # choose what training Dataset you want to USE (with augmentation or not??)
442 TrainingDataY = Y_Data_MIT_resampled_Training
443
444 return TrainingDataX, TrainingDataY, ValidationX, ValidationY,X_Test_MIT,Y_Test_MIT
445
446 # %% -----Plotting of data for inspection-----
447
448 # # %% Plotting of the Initial Data for Visualization
449 # fig = plt.figure(figsize=(15,4))
450
451 # for i in range(0,5):
452 #     plt.subplot(2,3,i + 1)
453 #     all_samples_indexes = np.where(Y_Data_MIT == i)[0]
454 #     rand_samples_indexes = np.random.randint(0, len(all_samples_indexes), 1)
455 #     rand_samples = X_Data_MIT[all_samples_indexes[5]] # plot the 5th of each class
456 #     plt.plot(rand_samples.transpose())
457 #     plt.ylim(-0.1, 1.3)
458
459 #     if i == 0 :
460 #         plt.ylabel('Norm Amplitude')
461 #     elif i == 3:
462 #         plt.ylabel('Norm Amplitude')
463 #     plt.title("Sample: " + classesMIT[i], loc='left', fontdict={'fontsize':10}, x=0.01, y=0.80)
464
465 # if saveData == 1:
466 #     fig.savefig('MIT_Data_Methods.png', format='png', dpi=800)
467
468 # %%----- Train and TEST Model on MIT DataSET -----
469
470 # lets load up the data for the model now
471 TrainingDataX, TrainingDataY, ValidationX, ValidationY,X_Test_MIT,Y_Test_MIT = data_curation_MIT()
472
473 n_classes = len(np.unique(TrainingDataY)) # how many classes are there in our data
474
475 # now pick which model we want to evaluate
476 if Model1_CNN:
477     print('Using Model 1')
478     ModelForTraining = make_model_1(n_classes) # Using the architecture CHANGE HERE
479
480 elif Model2_RNN:
481     print('Using Model 2')
482     ModelForTraining = make_model_2_RNN(n_classes) # Using the architecture CHANGE HERE
483
484 # for the RNN model we need to reshape the input a bit for it to work
485 TrainingDataX_RNN = np.reshape(TrainingDataX,(TrainingDataX.shape[0],TrainingDataX.shape[1],1))
486 ValidationX_RNN = np.reshape(ValidationX,(ValidationX.shape[0],ValidationX.shape[1],1))
487 TestDataX_RNN = np.reshape(X_Test_MIT,(X_Test_MIT.shape[0],X_Test_MIT.shape[1],1))
488
489 del TrainingDataX, TestSetX, ValidationX
490
491 TrainingDataX = TrainingDataX_RNN
492 ValidationX = ValidationX_RNN
493 X_Test_MIT = TestDataX_RNN
494
495 n_classes = len(np.unique(TrainingDataY))
496
497 elif Model3_OptimizationModel:
498     print('Using Model 3')
499     trials = Trials()
500
501 from os import path
502 pathFile = (r'E:\1_PAWELENSC_project\%s' % (ModelSaveName))
503 var = path.exists(pathFile)
504 if not var:
505
506     if Model3_OptimizationModel:
507         # Here is the setup for the optimized model, we set max_evals to 40 mainly as a hardware constraint as it takes quite a long time.
508         best_run, ModelOutput,return_space = optim.minimize(model=optimization_model,
509             data = data_curation_MIT,
510             algo=tpe.suggest,
511             max_evals=40,
512             trials=trials

```

```

524         return_space = True)
525
526
527     # saving it in this manner as to deal with issues of loading .h5 file after the optization. this seems to work now.
528     ModelSave = trials.results[np.argmax([r['loss'] for r in trials.results])]['model']
529     ModelSave.save(ModelSaveName)
530
531 else:
532
533     [ModelOutput, History]= NetworkModel(ModelForTraining, TrainingDataX, TrainingDataY, ValidationX, ValidationY, X_Test_MIT, Y_Test_MIT,
534                                           epochs, batch, patience, verbose, ModelSaveName, ComputeWeight=False, RNN=False, CNN= True, Optimzation = False )
535
536 else:
537
538     print('Loading preexisting model from folder .....')
539     ModelOutput = load_model(r'E:\1_PAWEL\ENSC_project\RESULTS\Model1_NoAUG\%s' % (ModelSaveName))
540
541     jsonName = ModelSaveName.replace('.h5', '')
542     hist_json_file = 'History_%s.json' %(jsonName)
543     with open(hist_json_file, 'r') as json_file:
544         History= json_file.read()
545
546
547     # this will not spit out the results, csv files etc into the folder to further look at
548     History = ModelOutput.history
549     CalculateModelMetrics(ModelOutput, History, X_Test_MIT, Y_Test_MIT, batch, Naming='MIT_resampled_Training_Weights.csv') #
550
551
552 plt.figure(figsize=(5, 5))
553 plt.plot(History.history['accuracy'])
554 plt.plot(History.history['val_accuracy'])
555 plt.title('model accuracy')
556 plt.ylabel('accuracy')
557 plt.xlabel('epoch')
558 plt.legend(['train', 'validation'], loc='upper left')
559 plt.show()
560
561 # summarize history for loss
562 # Plot non-normalized confusion matrix
563 plt.figure(figsize=(5, 5))
564 plt.plot(History.history['loss'])
565 plt.plot(History.history['val_loss'])
566 plt.title('model loss')
567 plt.ylabel('loss')
568 plt.xlabel('epoch')
569 plt.legend(['train', 'validation'], loc='upper left')
570 plt.show()
571
572 # %% Exploring Transfer Learning
573
574 # lets load in the other data to explore transfer learning
575 csv_PTB_Normal = pd.read_csv(r'E:\1_PAWEL\ENSC_project\heart_beatdata\ptbdb_normal.csv', header=None)
576 csv_PTB_AbNormal = pd.read_csv(r'E:\1_PAWEL\ENSC_project\heart_beatdata\ptbdb_abnormal.csv', header=None)
577 DataSet_PTB = pd.concat([csv_PTB_Normal, csv_PTB_AbNormal], axis=0)
578 Values_PTB = DataSet_PTB.values # get the values
579
580 X_Data_PTB = Values_PTB[:, :-1] # remove the last column which is the calssification
581 Y_Data_PTB = Values_PTB[:, -1] # keep the last column which will now be used for the Class Categoization
582
583
584 StratifyData = StratifiedShuffleSplit(n_splits=5, test_size=0.2, random_state=50)
585 # what we are trying to accomplish here is to split up the data such that the number of samples in each class is somewhat evenly distrubuted
586 # Since we already have a large class imbalance , this might be an issue if it is not addressed from the start. Lets split it up such that the
587 # test_size get 20% of the data (but catigorical samples evenly distrubuted)
588 StratifyData.get_n_splits(X_Data_PTB, Y_Data_PTB)
589
590 print(StratifyData)
591
592 for train_index, test_index in StratifyData.split(X_Data_PTB, Y_Data_PTB):
593     X_temp, X_Test_PTB = X_Data_PTB[train_index], X_Data_PTB[test_index]
594     y_temp, Y_Test_PTB = Y_Data_PTB[train_index], Y_Data_PTB[test_index]
595
596 uniqueValues, occurCountTrain = np.unique(y_temp, return_counts=True)
597 uniqueValues, occurCountValid = np.unique(Y_Test_PTB, return_counts=True)
598
599 print(occurCountTrain)
600 print(occurCountValid)
601
602 for train_index, test_index in StratifyData.split(X_temp, y_temp):
603
604     X_Train_PTB, X_Valid_PTB = X_temp[train_index], X_temp[test_index]
605     Y_Train_PTB, Y_Valid_PTB = y_temp[train_index], y_temp[test_index]
606
607 # # Plotting PTB database 1 sample from each cat
608
609 # plt.figure(figsize=(15,4))
610 # for i in range(0,2):
611 #     plt.subplot(1,2,i + 1)
612 #     all_samples_indexes = np.where(Y_Data_PTB == i)[0]
613 #     rand_samples_indexes = np.random.randint(0, len(all_samples_indexes), 1)
614 #     rand_samples = X_Data_PTB[all_samples_indexes[rand_samples_indexes]]
615 #     plt.plot(rand_samples.transpose())
616 #     plt.title("Sample: " + classePTB[i], loc="left", fontdict={'fontsize':10})
617
618
619 # change this line to change which model we want to evaluate
620 ModelOutputTransfer = ModelOutput # CHANGE THIS TO WHAT YOU WANT TO TEST
621
622 # lets build the new transfer learning layers ...
623 D1 = Dense(32)(ModelOutputTransfer.output)
624 D2 = Dense(32)(D1)
625
626 O = Dense(2, activation='softmax')(D2)
627 modelTransfer = Model(inputs=ModelOutputTransfer.input, outputs=O) # we have our new model
628
629 for layer in modelTransfer.layers[:-3]: # change this number to explore changing more layers in case of report i looked at -3, -5, -8
630     layer.trainable = False
631
632 for layer in modelTransfer.layers[-3:]: # change this number to explore changing more layers in case of report i looked at -3, -5, -8
633     layer.trainable = True

```

```
630 # we can use the function developed before and run it no on the new transfer learning model.
631 [modelTransferOutput, History]= NetworkModel(modelTransfer,X_Train_PTB,Y_Train_PTB,X_Valid_PTB,Y_Valid_PTB,X_Test_PTB,Y_Test_PTB,
632 epochs,batch,patience,verbose,ModelSaveName,ComputeWeight=False,RNN=False, CNN= True, Optimzation = False)

# lets output and look at how the transfer learning did on the novel dataset...
CalculateModelMetrics(modelTransferOutput,History,X_Test_PTB,Y_Test_PTB,batch, 'Transfer')
```



```

1 from tensorflow.keras.layers import Input, Conv1D, MaxPool1D, Activation, Add, Dense, Flatten
2 from tensorflow.keras.models import Model
3
4 def make_model_1(final_layer_size):
5     input_shape = (187, 1)
6     I = Input(input_shape)
7     C = Conv1D(filters=32, kernel_size=5, strides=1)(I)
8
9     C11 = Conv1D(filters=32, kernel_size=5, strides=1, padding='same')(C)
10    R11 = Activation(activation='relu')(C11)
11    C12 = Conv1D(filters=32, kernel_size=5, strides=1, padding='same')(R11)
12    A11 = Add()([C12, C])
13    R12 = Activation(activation='relu')(A11)
14    M11 = MaxPool1D(pool_size=5, strides=2)(R12)
15
16    C21 = Conv1D(filters=32, kernel_size=5, strides=1, padding='same')(M11)
17    R21 = Activation(activation='relu')(C21)
18    C22 = Conv1D(filters=32, kernel_size=5, strides=1, padding='same')(R21)
19    A21 = Add()([C22, M11])
20    R22 = Activation(activation='relu')(A21)
21    M21 = MaxPool1D(pool_size=5, strides=2)(R22)
22
23    C31 = Conv1D(filters=32, kernel_size=5, strides=1, padding='same')(M21)
24    R31 = Activation(activation='relu')(C31)
25    C32 = Conv1D(filters=32, kernel_size=5, strides=1, padding='same')(R31)
26    A31 = Add()([C32, M21])
27    R32 = Activation(activation='relu')(A31)
28    M31 = MaxPool1D(pool_size=5, strides=2)(R32)
29
30    C41 = Conv1D(filters=32, kernel_size=5, strides=1, padding='same')(M31)
31    R41 = Activation(activation='relu')(C41)
32    C42 = Conv1D(filters=32, kernel_size=5, strides=1, padding='same')(R41)
33    A41 = Add()([C42, M31])
34    R42 = Activation(activation='relu')(A41)
35    M41 = MaxPool1D(pool_size=5, strides=2)(R42)
36
37    C51 = Conv1D(filters=32, kernel_size=5, strides=1, padding='same')(M41)
38    R51 = Activation(activation='relu')(C51)
39    C52 = Conv1D(filters=32, kernel_size=5, strides=1, padding='same')(R51)
40    A51 = Add()([M41, C52])
41    R52 = Activation(activation='relu')(A51)
42    M51 = MaxPool1D(pool_size=5, strides=2)(R52)
43
44    F1 = Flatten()(M51)
45
46    D1 = Dense(32)(F1)
47    R6 = Activation(activation='relu')(D1)
48    D2 = Dense(32)(R6)
49    D3 = Dense(final_layer_size)(D2)
50    O = Activation(activation='softmax')(D3)
51
52    return Model(inputs=I, outputs=O)

```

```
1 from tensorflow.keras.layers import Input,Dropout, LSTM, Bidirectional,BatchNormalization, Dense
2 from tensorflow.keras.models import Sequential
3
4
5 def make_model_2_RNN(final_layer_size):
6
7
8     Model = Sequential()
9     Model.add(Dense(32, input_shape=(187,1)))
10    Model.add(Bidirectional(LSTM(64,input_shape = (187,1))))
11    Model.add(Dropout(rate =0.25))
12    Model.add(BatchNormalization())
13    Model.add(Dense(32, activation='relu'))
14    Model.add(Dropout(rate =0.25))
15    Model.add(BatchNormalization())
16    Model.add(Dense(final_layer_size,activation = 'sigmoid'))
17
18
19
20    return Model
```